

# Carrier Recovery

Joe Leacy

May 13, 2026

## 1 Introduction

## 2 Frequency Recovery

### 2.1 Isolating CFO

All frequency recovery algorithms first require that the phase of the data be removed before estimation. For QPSK symbols  $c[i] = e^{j\frac{\pi}{4} + m[i]\frac{\pi}{2}}$ , this can be done by raising each received symbol  $x[i] = c[i]e^{j(2\pi f_d i + \theta[i])} + n[i]$  to the fourth power to give

$$z[i] = x^4[i] = e^{j(8\pi f_d i + 4\theta[i])} + n'[i] \quad (1)$$

where  $\theta$  is phase noise,  $f_d$  is the CFO, and  $n$  is additive noise. This allows all data symbols to be used in CFO estimation, but amplifies phase and additive noise which degrades the estimation quality. The fourth power operation also requires three complex multiplications. An alternative is to use pilot symbols where the data phase is known beforehand, giving

$$z[i] = x[i]c^*[i] = e^{j(2\pi f_d i + \theta[i])} + n[i], \quad (2)$$

which is more robust to noise and less computationally expensive. Some algorithms only require  $\arg(z[i])$  which is efficiently found using CORDIC as

$$\arg(z[i]) = \{4 \arg(x[i])\}|_{-\pi}^{\pi} \quad (3)$$

for blind operation or

$$\arg(z[i]) = \{\arg(x[i]) - \arg(c[i])\}|_{-\pi}^{\pi} \quad (4)$$

for data-aided operation where the wrap operation can be done with no additional cost by aligning the wrap range to the overflow range of the fixed-point value used to store the phase.

### 2.2 Rife and Boorstyn

The value of the CFO will appear as the strongest peak in the frequency spectrum of  $z[i], z[i+1], \dots, z[i+L]$  where  $L$  is the observation length. The Rife and Boorstyn (R&B) algorithm first performs an FFT on the observed sequence and finds the index  $k$  corresponding to the strongest peak. This coarse estimation is refined with a parabolic interpolation on the complex FFT values  $X_{k-1}, X_k, X_{k+1}$  [3] to give

$$\delta = -\Re \left[ \frac{X_{k+1} - X_{k-1}}{2X_k - X_{k-1} - X_{k+1}} \right]. \quad (5)$$

The updated index  $k' = k + \delta$ , which is then scaled by the sampling rate to find the true CFO. The CPON specification allows for 11 training symbols at the start of each subframe. The resolution in  $k$  is determined by the FFT size  $N_{FFT}$  as  $\epsilon(k) = 1/2N_{FFT}$ . For data-aided operation, only the 11 training symbols are available. This would give a resolution of  $\epsilon(k) \approx 0.05$ . In terms of the estimated frequency  $\hat{f}_d$  and symbol rate  $R_s$ ,  $\epsilon(\hat{f}_d) \approx 0.05R_s$ . This would likely be too large for the downstream phase recovery to correct, so a larger FFT can be used with zeros padded at the end of the observed sequence. The normalised mean squared-error (MSE) in the estimated CFO is shown in fig. 1 for data-aided operation with 11 training symbols. A CFO estimate is computed for both polarizations and averaged to reduce the MSE. The modified Cramer-Rao bound (MCRB) [2] provides a lower bound on the achievable MSE and is given by

$$\text{MCRB} = \frac{3}{2\pi^2 L^3 \times \text{SNR}}. \quad (6)$$

R&B achieves this bound above a threshold SNR of  $\approx 4$  dB, where the threshold SNR is defined as the SNR below which the estimator begins to deviate from the MCRB. Below this, the noise has a significant contribution to the FFT values and can lead to an incorrect index being chosen.




Figure 1: Normalised MSE of the R&B algorithm with FFT size  $N_{FFT} = 512$  compared to the MCRB for  $L = 11$  training symbols.

Data-aided operation is constrained by the length of the training sequence, but the accuracy of the R&B algorithm improves with longer observation lengths, as shown in fig. 2 for non-data-aided operation. The SNR threshold at which this algorithm begins to deviate from the MCRB is initially higher with data-aided operation due to the amplification of noise, but decreases for longer observation lengths.




Figure 2: Normalised MSE of the R&B algorithm for non-data-aided operation and different observation lengths.

### 2.3 Differential Phase + Kay Estimator

The R&B algorithm exhibits good performance even at low SNRs but is computationally expensive. The FFT requires  $\frac{N_{FFT}}{2} \log_2(N_{FFT})$  complex multiplications and  $N_{FFT} \log_2(N_{FFT})$  complex additions. The fine search from eq. (5) also requires a complex division. A cheaper approach would involve computing the phase difference

$$\arg(z[i]) - \arg(z[i-1]) = 2\pi f_d T_s + \theta[i] - \theta[i-1] + \eta[i] \quad (7)$$

over several symbols and averaging to remove the effect of  $\theta$ . Estimating  $f_d$  in this manner is vulnerable to the effect of additive noise on the phase  $\eta[i]$  so exhibits poor accuracy. This can be improved with a second pass using the Kay estimator. The Kay estimator uses a least-squares approach, computing the estimated CFO as

$$\hat{f}_d = \frac{R_s}{2\pi} \sum_{k=1}^L \frac{6k(L-k)}{L(L^2-1)} \arg(z[k]z^*[k-1]). \quad (8)$$

where  $\arg(z[k]z^*[k-1])$  computes the phase difference from eq. (7) and  $\frac{6k(L-k)}{L(L^2-1)}$  is a weighting term that prioritises the contribution of phase differences in the middle of the observed sequence. This estimator is able to achieve the MCRB but the least-squares approach considers phase differences of non-adjacent symbols, making it vulnerable to phase wrapping that distorts the estimate when the CFO is high enough. This is demonstrated by considering the phase of symbols at either end of the observation sequence. The phase drift between these symbols  $\Delta\phi = 2\pi f_d L T_s$ . If  $|\Delta\phi| > 2\pi$  then a phase wrap is guaranteed to occur, giving an estimated phase drift  $\Delta\hat{\phi} = \Delta\phi \pm 2\pi$  and an estimated CFO  $\hat{f}_d = f_d \pm \frac{1}{L T_s}$ . Thus, a first-pass CFO correction is needed to bring the residual CFO within range for the Kay estimator. The performance of the Differential Kay (DK) estimator for data-aided operation is shown in fig. 3. The MCRB is achieved above a threshold SNR of  $\approx 7$  dB. This is higher than for R&B but at a reduced computational cost. Unlike the R&B algorithm, the threshold SNR of the DK estimator does not depend on observation length, as shown in fig. 4.




Figure 3: Normalised MSE of frequency estimation with the DK estimator for  $L = 11$  training symbols.




Figure 4: Normalised MSE of the DK estimator with blind operation for different observation lengths.

## 2.4 Computational Complexity

Here the computational complexity of both frequency recovery algorithms is compared in terms of real multiplications and real additions. CORDIC is used to find phases and is assumed to cost the same as a real multiplication. Real division will also be assumed to cost the same as multiplication as it is only performed once (this would be a suitable approximation if look-up tables are used to find reciprocals at some cost to accuracy). All analysis is for data-aided operation, avoiding the fourth power operation in eq. (1), and applies to a single polarization. If averaging is done over both polarizations, then all operations are doubled.

**R&B** The R&B algorithm requires an FFT followed by  $N_{FFT}$  comparisons of the magnitudes in the coarse search. Many of the FFT entries are 0 due to padding, which reduces the total number of operations needed. The interpolation requires the real part of a complex division

$$\Re \left[ \frac{a + bj}{c + dj} \right] = \frac{ac + bd}{c^2 + d^2}. \quad (9)$$

The total complexity is summarised in table 1 with the complexity for blind operation marked with (\*).

| Operation                       | RM                                                  | RA                                                  |
|---------------------------------|-----------------------------------------------------|-----------------------------------------------------|
| Forming $z[i], \dots, z[i + L]$ | $4L, 12L^*$                                         | $3L, 9L^*$                                          |
| FFT                             | $2L \log_2(\frac{N_{FFT}}{L}) + 2N_{FFT} \log_2(L)$ | $3L \log_2(\frac{N_{FFT}}{L}) + 3N_{FFT} \log_2(L)$ |
| Search                          | $2N_{FFT}$                                          | $2N_{FFT}$                                          |
| Interpolation                   | 5                                                   | 4                                                   |

Table 1: Computational complexity of estimating CFO with the R&B algorithm.

**DK** The DK estimator has lower complexity, only requiring CORDIC, real additions, and real multiplications. The fixed multiplications by the weights of the Kay estimator can be optimised into a low-energy shift-add circuit, but this is ignored for simplicity. The first-pass CFO correction involves forming the phase corrections of each training symbol as  $\phi[i] = \hat{f}_d i$ , then applying them with CORDIC.

| Operation                                   | RM   | RA      |
|---------------------------------------------|------|---------|
| Forming $\arg(z[i]), \dots, \arg(z[i + L])$ | $2L$ | $L$     |
| Differential Phase                          | 1    | $L - 1$ |
| First Pass Correction                       | $2L$ | 0       |
| Forming $\arg(z[i]), \dots, \arg(z[i + L])$ | $2L$ | $L$     |
| Kay                                         | $L$  | $L - 1$ |

Table 2: Computational complexity of estimating CFO with the DK estimator.

## 2.5 Optimising Bit Width

These frequency recovery algorithms compute an estimate of the normalised CFO  $f_d/R_s$ . For  $n$  bits of precision, the smallest normalised CFO that can be estimated is  $2^{-n}$  when  $0 \leq |f_d/R_s| \leq 1$  is the estimation range. In reality a much smaller range of CFOs needs to be estimated. If the maximum permitted value of  $|f_d/R_s| = 2^{-d}$ ,  $d$  bits of precision that would otherwise be redundant can be utilised by scaling each estimate by  $2^d$  during calculation then inserting  $d$  bits in the final answer to recover the true normalised CFO. This improves the resolution to  $2^{-(n+d)}$  without any increase in energy consumption.

## 3 Phase Recovery

### 3.1 Feed-Forward vs PLL Algorithms

One of the first techniques investigated for phase recovery was a digital phase-locked loop (PLL) to incrementally correct the phase of each symbol with feedback. This dramatically simplifies the phase estimation and correction due to the use of small angle approximations. The CPON spec sets the symbol rate at 30.5 GBd. With an oversampling rate of 8/7 the total samples per second would be 35 GSa/s. This is beyond the achievable clock speed of CMOS, so symbols must be processed in blocks of e.g., 32 symbols with a clock speed of 1.1 GHz. This approach hinders the effectiveness of a PLL, as updates can only be made every 32 symbols. For this reason, feed-forward algorithms that can operate on blocks of data were used instead.

### 3.2 Viterbi & Viterbi

The Viterbi & Viterbi algorithm is well suited to QPSK and removes data phase with the fourth power operation of eq. (1). It then applies a maximum-likelihood filter  $w_{ML}[i]$  to estimate the phase error due to Wiener phase noise from the combined laser linewidth and any residual CFO as

$$\hat{\theta}_{ML}[i] = \frac{1}{4} \arg \left( \sum_{i=1}^N w_{ML}^*[i] z[i] \right) - \frac{\pi}{4}. \quad (10)$$

A phase unwrapper is then applied before the phase error is corrected. This is readily adapted to block-based processing. If blocks of 32 symbols are used again, the ML filter can be applied on all symbols in the block and the estimated phase error applied to the entire block. This is only valid if the phase difference between the start and end of the block is small, meaning the laser linewidth cannot be too large and the preceding frequency recovery algorithm must correct most of the CFO. By increasing the time between phase estimates, block processing also increases the chance of cycle-slips in the phase unwrapper. The CPON spec supplies 1 pilot symbol in every 32-symbol block that can be used to correct this [1]. The correction estimates the number of  $\pi/2$  phase jumps  $n$  that occur between the unwrapped phase of a data symbol  $\hat{\theta}_{PU}$  and its corresponding pilot, then adjusts the unwrapped estimate as  $\hat{\theta}'_{PU} = \hat{\theta}_{PU} - n\frac{\pi}{2}$ .

#### 3.2.1 Pilots-Only Correction

While Viterbi & Viterbi is commonly used for QPSK, a simpler approach is to compute the phase error from the pilot symbol using eq. (2) and apply this to the whole block. This also relies on a small phase change across the block and is more vulnerable to additive noise. However, QPSK allows for  $\pm 45^\circ$  error in the phase estimate before an incorrect decision is made, which makes it robust to the drop in accuracy associated with this approach.

### 3.3 Computational Complexity

**Viterbi & Viterbi** For a block-based approach, the length of the filter will be the block length. Phases are again found with CORDIC. The cost of the phase estimate for each block is shown in table 3.

| Operation                       | RM       | RA                  |
|---------------------------------|----------|---------------------|
| Forming $z[i], \dots, z[i + N]$ | $12N$    | $9N$                |
| Computing $\hat{\theta}_{ML}$   | $4N + 1$ | $3N + 2(N - 1) + 1$ |

Table 3: Computational complexity of block-based Viterbi & Viterbi.

**Pilots-Only** Only phases are needed, so the entire estimation can be performed with CORDIC. The pilot phase is also known beforehand, so does not contribute to the computational cost. The total cost of the phase estimate is shown in table 4.

| Operation                           | RM | RA |
|-------------------------------------|----|----|
| Forming $\arg(z[i]) = \hat{\theta}$ | 1  | 1  |

Table 4: Computational complexity of Pilots-Only correction.

## 4 Amortisation

The estimated CFO can be applied to the number of symbols that it can be assumed to remain constant over, dividing the operation counts in table 1 and table 2 by this number gives the operations per symbol. The phase recovery stage operates on blocks of 32 symbols, so the operations per symbol is found by dividing the total operations from table 3 and table 4 by the block length.

## 5 Methodology

## 6 Results & Discussion

### 6.1 Maximum CFO

The theoretical limit on the maximum CFO that can be corrected is given by Nyquist as  $0.5R_s$ . From eq. (1) we see that blind operation reduces this limit by a factor of 4. Both the R&B and DK estimators accurately estimate the CFO up to the Nyquist limit as seen in fig. 5. Blind operation aliases as expected at  $f/R_s = 0.125$ .




Figure 5: Normalised MSE in CFO estimate ( $\text{MSE}^2/R_s^2$ ) vs Normalised CFO ( $f/R_s$ )

## 6.2 Blind vs Data-Aided Operation

The decrease in threshold SNR for longer observation lengths suggests that the performance of the R&B algorithm will improve with blind observation length. This is not the case for the DK estimator, where increasing blind observation length has no effect on threshold SNR. The effect of observation length on FEC SNR is shown for both algorithms in fig. 6, where blind R&B shows a clear increase in performance over the blind DK estimator for longer observations and improves on data-aided R&B after  $\approx 50$  symbols. The DK estimator in blind operation is thus not viable for frequency recovery but blind R&B can provide the best performance at a higher computational cost.




Figure 6: FEC SNR vs blind observation length for the R&B and DK estimators. CFO = 2 GHz, laser linewidth = 1 MHz



### 6.3 Combined Performance

The energy consumption of any combination of frequency and phase recovery algorithms can be optimised by varying the precision of the fixed point arithmetic in each stage. Blind operation also allows for optimisation over the observation length. The lowest receiver energy per bit for a given FEC SNR is shown for each combination in fig. 7. Here the CFO is assumed to stay constant over each subframe of length 3712 symbols so the amortisation of each stage means that the choice of phase recovery algorithm dominates the receiver energy. R&B can also use less energy in blind operation when allowing fixed point precision to vary, as the more accurate CFO estimate can use fewer bits of precision to achieve the same FEC SNR.




Figure 7: Lowest receiver energy for each FEC SNR achieved.

### 6.4 System Energy

If carrier recovery limits the required SNR at the receiver, the FEC SNR and the receiver energy of each carrier recovery combination will determine the system energy for acceptable operation. The minimum energy used by the transmitter and carrier recovery block to achieve the FEC limit, along with the SNR at which it is achieved, is shown in table 5 for both amplifier noise and shot noise limited links, and for  $K = 1, 10$ , and 100 ONUs per OLT.

| Algorithm pair       | SNR (dB) | $K = 1$     |              | $K = 10$    |             | $K = 100$   |             |
|----------------------|----------|-------------|--------------|-------------|-------------|-------------|-------------|
|                      |          | Amp.        | Shot         | Amp.        | Shot        | Amp.        | Shot        |
| R&B (DA) + V&V       | 17.71    | 994         | 981          | 9820        | 9810        | 98100       | 98100       |
| R&B (DA) + Pilots    | 10.11    | 29.6        | 27.2         | 275         | 272         | 2720        | 2720        |
| DK (DA) + V&V        | 17.74    | 968         | 954          | 9550        | 9540        | 95400       | 95400       |
| DK (DA) + Pilots     | 10.73    | <b>3.69</b> | <b>0.940</b> | <b>11.8</b> | <b>9.40</b> | <b>93.0</b> | <b>94.1</b> |
| R&B (blind) + V&V    | 17.65    | 1010        | 992          | 9930        | 9920        | 99200       | 99200       |
| R&B (blind) + Pilots | 10.30    | 41.2        | 38.7         | 390         | 387         | 3870        | 3870        |

Table 5: Minimum combined energy (fJ/bit) for each carrier recovery combination under amplifier noise and shot noise limited links for  $K$  ONUs per OLT.

## References

- [1] Haiquan Cheng et al. “Pilot-symbols-aided cycle slip mitigation for DP-16QAM optical communication systems”. In: *Optics Express* 21.19 (2013), pp. 22166–22172. DOI: 10.1364/OE.21.022166.
- [2] A.N. D’Andrea, Umberto Mengali, and Ruggero Reggiannini. “The Modified Cramer-Rao Bound and Its Applications to Synchronization Problems”. In: *Communications, IEEE Transactions on* 42 (Mar. 1994), pp. 1391–1399. DOI: 10.1109/TCOMM.1994.580247.
- [3] Eric Jacobsen and Peter Kootsookos. “Fast, Accurate Frequency Estimators”. In: *IEEE Signal Processing Magazine* 24.3 (May 2007), pp. 123–125. DOI: 10.1109/MSP.2007.361611.