

Energy-Efficient 200 Gbit/s Transceivers for Optical Access Networks

Joe Leacy

May 13, 2026

Contents

1	Estimating Energy Consumption	2
1.1	Introduction	2
1.2	Receiver Energy	2
1.2.1	Dependence on Word Length	3
1.2.2	Amortisation	3
1.2.3	Total Energy and Power	3
1.3	Transmitter Energy	3
1.3.1	Minimum Required SNR	4
1.3.2	Noise	4
1.4	System Energy	5
2	Carrier Recovery	6
2.1	Introduction	6
2.2	Frequency Recovery	6
2.2.1	Isolating CFO	6
2.2.2	Rife and Boorstyn	6
2.2.3	Differential Phase + Kay Estimator	8
2.2.4	Computational Complexity	9
2.2.5	Optimising Bit Width	10
2.3	Phase Recovery	11
2.3.1	Feed-Forward vs PLL Algorithms	11
2.3.2	Viterbi & Viterbi	11
2.3.3	Computational Complexity	11
2.4	Methodology	12
2.5	Results & Discussion	12
2.5.1	Maximum CFO	12
2.5.2	Blind vs Data-Aided Operation	13
2.5.3	System Energy	13

Chapter 1

Estimating Energy Consumption

1.1 Introduction

The energy efficiency of a digital signal processing (DSP) technique used in a coherent receiver cannot be judged from the receiver alone: a simpler implementation that saves receiver energy typically tolerates a lower SNR, which the transmitter must compensate for with extra launch power. A like-for-like comparison between architectures must therefore account for both contributions.

This section develops such a framework. The receiver contribution is modelled from operation counts of real additions and multiplications, with per-operation energies that depend on the fixed-point word length, allowing implementations to be compared without a hardware build. The transmitter contribution is derived from the minimum SNR each implementation requires to meet the pre-FEC BER threshold defined in the coherent PON (CPON) standard [1] under investigation. The two are combined for a downstream PON to give the per-bit system energy

1.2 Receiver Energy

The energy consumption of the receiver can be estimated from the number of real multiplications N_M and real additions N_A it performs to process each symbol. The true cost of these operations depends on the hardware implementation. Here the reference values shown in table 1.1 are used to allow for comparison between different DSP algorithms. An accurate estimate of their actual energy consumption would require analysis of their ASIC implementation for a modern process node.

Table 1.1: Energy cost of arithmetic operations for different bit widths in 45 nm CMOS [4].

n (bits)	Add (pJ)	Multiply (pJ)
8	0.03	0.2
32	0.1	3.1

1.2.1 Dependence on Word Length

A straightforward adder implementation performs n bit-wise additions on words of length n , so its energy scales linearly with n . Multiplication accumulates n partial products of length n and so scales as n^2 . These physical scaling laws are imposed *a priori* and the proportionality constants calibrated by least-squares fit through the origin to the two data points of table 1.1, giving

$$E_A(n) = 3.16 n \text{ fJ}, \quad E_M(n) = 3.03 n^2 \text{ fJ}. \quad (1.1)$$

1.2.2 Amortisation

Some DSP algorithms perform computations that are reused over multiple symbols. For example, carrier frequency offset (CFO) varies much slower than the symbol rate so can be estimated less frequently. For a computation that requires N_M multiplications and N_A additions and is used by S symbols, the operations per symbol would be N_M/S and N_A/S .

1.2.3 Total Energy and Power

The energy per symbol for the entire receiver is given by

$$E_{\text{rx},s} = \sum_{k=1}^L 2 [N_{A,k} E_A(n_k) + N_{M,k} E_M(n_k)], \quad (1.2)$$

where the receiver is decomposed into L blocks that use different word lengths n_k according to their required precision, and $N_{A,k}$, $N_{M,k}$ are the addition and multiplication counts per symbol. The factor of 2 is for both polarizations. Blocks operating on the oversampled signal at rate $f > R_s$ have their per-sample operation counts scaled by the oversampling ratio f/R_s . The receiver power is then

$$P_{\text{rx}} = R_s E_{\text{rx},s}, \quad (1.3)$$

and converting to energy per bit allows fair comparison across modulation schemes and symbol rates,

$$E_{\text{rx}} = \frac{E_{\text{rx},s}}{\log_2(M)}, \quad (1.4)$$

where M is the modulation order. Complex multiplications and additions are needed in many DSP algorithms. 4 real multiplications and 2 real additions are assumed for each complex multiplication, and 2 real additions for complex addition.

1.3 Transmitter Energy

The downstream link considered here consists of a single OLT transmitting through a 1:K passive splitter to K ONUs, as shown in fig. 1.1. The OLT power is divided equally between branches, so each ONU receives a $1/K$ share. This is a simplification of a typical PON where there may be several stages of passive splitters placed at different locations along the fibre, in which case the noise characteristics of the fibre between splitter stages must be analysed independently.




Figure 1.1: Downstream PON link: one OLT transmitting to K ONUs.

1.3.1 Minimum Required SNR

The CPON standard specifies a maximum pre-FEC bit-error rate (BER) of 10^{-2} that must not be exceeded for the forward error-correction (FEC) scheme to function correctly. Each receiver implementation therefore has a minimum required SNR to meet this FEC threshold. This SNR, together with the noise characteristics of the optical network, determines the minimum transmitter power.

1.3.2 Noise

Shot Noise

The relationship between SNR and transmitter power P_{tx} depends on the dominant source of noise in the network. Optical access networks typically operate over short lengths at low power, so amplifiers can be omitted and non-linear noise neglected, leaving shot noise as the dominant impairment. The shot noise power at the receiver is $2h\nu B_{\text{eff}}$, where B_{eff} is the bandwidth of the photodetector; this can be matched to the bandwidth of the received signal after the 1:K splitter in fig. 1.1 with a bandpass filter. The received power after attenuation and the splitter (assuming equal power per channel) is $P_{\text{tx}} \exp(-\alpha L)/K$, giving a per-ONU SNR of

$$\text{SNR} = \frac{P_{\text{tx}} \exp(-\alpha L)}{2Kh\nu B_{\text{eff}}}, \quad (1.5)$$

where α is the fibre loss in nepers/km, and L is the fibre length in km. Rearranging gives the transmitted power as

$$P_{\text{tx}} = 2Kh\nu B_{\text{eff}} \exp(\alpha L) \cdot \text{SNR}, \quad (1.6)$$

which is proportional to K when channels transmit equal power. The energy per bit transmitted is then

$$E_{\text{tx}} = \frac{P_{\text{tx}}}{K R_s \log_2(M)}. \quad (1.7)$$

Amplifier and Non-Linear Noise

Long-reach access networks require EDFAs in the trunk fibre, in which case amplifier and non-linear noise can no longer be neglected. Each EDFA with gain G and noise figure NF (dB) operating over a full WDM bandwidth B contributes an ASE noise power

$$P_{\text{ASE}} = N_{\text{ASE}} B = 10^{\text{NF}/10} h\nu (G - 1) B, \quad (1.8)$$

and assuming each amplifier exactly compensates the loss of its preceding span, every span re-launches with the same input power P_{tx} . The Gaussian Noise model [6] then gives the per-span non-linear interference (NLI) power as

$$P_{\text{NLI}} = \frac{C_{\text{NLI}} P_{\text{tx}}^3}{B^2}, \quad (1.9)$$

where $C_{\text{NLI}} = \frac{8\gamma^2 L_{\text{eff}}^2 \alpha}{27\pi |\beta_2|} \ln(\frac{|\beta_2|}{\alpha} \pi^2 B^2)$ depends on the fibre and link parameters.

For a link of n identical spans, the ASE and NLI contributions from each span accumulate, giving a total noise power $n(P_{\text{ASE}} + P_{\text{NLI}})$. With the 1:K passive splitter placed at the end of the trunk, both signal and accumulated noise are attenuated equally between the OLT and each ONU, so the splitter ratio cancels in the per-ONU SNR:

$$\text{SNR} = \frac{P_{\text{tx}}}{n \left(N_{\text{ASE}} B + \frac{C_{\text{NLI}} P_{\text{tx}}^3}{B^2} \right)}, \quad (1.10)$$

where P_{tx} is both the OLT transmit power and (under the equal-compensation assumption) the launch power into each span. The required transmit powers for each implementation are solved for numerically from eq. (1.10), with the transmitter energies following from eq. (1.7). Here the transmitted power is not proportional to K for a given SNR even if channels transmit equal power due to the dependence of C_{NLI} on the bandwidth.

1.4 System Energy

The total energy consumed per bit for network is given by

$$E_{\text{sys}} = \frac{1}{\eta} E_{\text{tx}} + E_{\text{rx}}, \quad (1.11)$$

where η is the wall-plug efficiency of the OLT laser. In the shot noise regime, E_{sys} is independent of K , but will vary when amplifier and non-linear noise dominate.

Chapter 2

Carrier Recovery

2.1 Introduction

2.2 Frequency Recovery

2.2.1 Isolating CFO

All frequency recovery algorithms first require that the phase of the data be removed before estimation. For QPSK symbols $c[i] = e^{j\frac{\pi}{4} + m[i]\frac{\pi}{2}}$, this can be done by raising each received symbol $x[i] = c[i]e^{j(2\pi f_d i + \theta[i])} + n[i]$ to the fourth power to give

$$z[i] = x^4[i] = e^{j(8\pi f_d i + 4\theta[i])} + n'[i] \quad (2.1)$$

where θ is phase noise, f_d is the CFO, and n is additive noise. This allows all data symbols to be used in CFO estimation, but amplifies phase and additive noise which degrades the estimation quality. The fourth power operation also requires three complex multiplications. An alternative is to use pilot symbols where the data phase is known beforehand, giving

$$z[i] = x[i]c^*[i] = e^{j(2\pi f_d i + \theta[i])} + n[i], \quad (2.2)$$

which is more robust to noise and less computationally expensive. Some algorithms only require $\arg(z[i])$ which is efficiently found using CORDIC as

$$\arg(z[i]) = \{4 \arg(x[i])\}|_{-\pi}^{\pi} \quad (2.3)$$

for blind operation or

$$\arg(z[i]) = \{\arg(x[i]) - \arg(c[i])\}|_{-\pi}^{\pi} \quad (2.4)$$

for data-aided operation where the wrap operation can be done with no additional cost by aligning the wrap range to the overflow range of the fixed-point value used to store the phase.

2.2.2 Rife and Boorstyn

The value of the CFO will appear as the strongest peak in the frequency spectrum of $z[i], z[i+1], \dots, z[i+L]$ where L is the observation length. The Rife and Boorstyn (R&B) algorithm first performs an FFT on the observed sequence and finds the index k

corresponding to the strongest peak. This coarse estimation is refined with a parabolic interpolation on the complex FFT values X_{k-1}, X_k, X_{k+1} [5] to give

$$\delta = -\Re \left[\frac{X_{k+1} - X_{k-1}}{2X_k - X_{k-1} - X_{k+1}} \right]. \quad (2.5)$$

The updated index $k' = k + \delta$, which is then scaled by the sampling rate to find the true CFO. The CPON specification allows for 11 training symbols at the start of each subframe. The resolution in k is determined by the FFT size N_{FFT} as $\epsilon(k) = 1/2N_{FFT}$. For data-aided operation, only the 11 training symbols are available. This would give a resolution of $\epsilon(k) \approx 0.05$. In terms of the estimated frequency $\hat{f}_d$ and symbol rate R_s , $\epsilon(\hat{f}_d) \approx 0.05R_s$. This would likely be too large for the downstream phase recovery to correct, so a larger FFT can be used with zeros padded at the end of the observed sequence. The normalised mean squared-error (MSE) in the estimated CFO is shown in fig. 2.1 for data-aided operation with 11 training symbols. A CFO estimate is computed for both polarizations and averaged to reduce the MSE. The modified Cramer-Rao bound (MCRB) [3] provides a lower bound on the achievable MSE and is given by

$$\text{MCRB} = \frac{3}{2\pi^2 L^3 \times \text{SNR}}. \quad (2.6)$$

R&B achieves this bound above a threshold SNR of ≈ 4 dB, where the threshold SNR is defined as the SNR below which the estimator begins to deviate from the MCRB. Below this, the noise has a significant contribution to the FFT values and can lead to an incorrect index being chosen.




Figure 2.1: Normalised MSE of the R&B algorithm with FFT size $N_{FFT} = 512$ compared to the MCRB for $L = 11$ training symbols.

Data-aided operation is constrained by the length of the training sequence, but the accuracy of the R&B algorithm improves with longer observation lengths, as shown in fig. 2.2 for non-data-aided operation. The SNR threshold at which this algorithm begins to deviate from the MCRB is initially higher with data-aided operation due to the amplification of noise, but decreases for longer observation lengths.




Figure 2.2: Normalised MSE of the R&B algorithm for non-data-aided operation and different observation lengths.

2.2.3 Differential Phase + Kay Estimator

The R&B algorithm exhibits good performance even at low SNRs but is computationally expensive. The FFT requires $\frac{N_{FFT}}{2} \log_2(N_{FFT})$ complex multiplications and $N_{FFT} \log_2(N_{FFT})$ complex additions. The fine search from eq. (2.5) also requires a complex division. A cheaper approach would involve computing the phase difference

$$\arg(z[i]) - \arg(z[i-1]) = 2\pi f_d T_s + \theta[i] - \theta[i-1] + \eta[i] \quad (2.7)$$

over several symbols and averaging to remove the effect of θ . Estimating f_d in this manner is vulnerable to the effect of additive noise on the phase $\eta[i]$ so exhibits poor accuracy. This can be improved with a second pass using the Kay estimator. The Kay estimator uses a least-squares approach, computing the estimated CFO as

$$\hat{f}_d = \frac{R_s}{2\pi} \sum_{k=1}^L \frac{6k(L-k)}{L(L^2-1)} \arg(z[k]z^*[k-1]). \quad (2.8)$$

where $\arg(z[k]z^*[k-1])$ computes the phase difference from eq. (2.7) and $\frac{6k(L-k)}{L(L^2-1)}$ is a weighting term that prioritises the contribution of phase differences in the middle of the observed sequence. This estimator is able to achieve the MCRB but the least-squares approach considers phase differences of non-adjacent symbols, making it vulnerable to phase wrapping that distorts the estimate when the CFO is high enough. This is demonstrated by considering the phase of symbols at either end of the observation sequence. The phase drift between these symbols $\Delta\phi = 2\pi f_d L T_s$. If $|\Delta\phi| > 2\pi$ then a phase wrap is guaranteed to occur, giving an estimated phase drift $\Delta\hat{\phi} = \Delta\phi \pm 2\pi$ and an estimated CFO $\hat{f}_d = f_d \pm \frac{1}{L T_s}$. Thus, a first-pass CFO correction is needed to bring the residual CFO within range for the Kay estimator. The performance of the Differential Kay (DK) estimator for data-aided operation is shown in fig. 2.3. The MCRB is achieved above a threshold SNR of ≈ 7 dB. This is higher than for R&B but at a reduced computational cost. Unlike the R&B algorithm, the threshold SNR of the DK estimator does not depend on observation length, as shown in fig. 2.4.




Figure 2.3: Normalised MSE of frequency estimation with the DK estimator for $L = 11$ training symbols.




Figure 2.4: Normalised MSE of the DK estimator with blind operation for different observation lengths.

2.2.4 Computational Complexity

Here the computational complexity of both frequency recovery algorithms is compared in terms of real multiplications and real additions. CORDIC is used to find phases and is assumed to cost the same as a real multiplication. Real division will also be assumed to cost the same as multiplication as it is only performed once (this would be a suitable approximation if look-up tables are used to find reciprocals at some cost to accuracy). All analysis is for data-aided operation, avoiding the fourth power operation in eq. (2.1), and applies to a single polarization. If averaging is done over both polarizations, then all operations are doubled.

R&B The R&B algorithm requires an FFT followed by N_{FFT} comparisons of the magnitudes in the coarse search. Many of the FFT entries are 0 due to padding, which reduces the total number of operations needed. The interpolation requires the real part of a complex division

$$\Re \left[\frac{a + bj}{c + dj} \right] = \frac{ac + bd}{c^2 + d^2}. \quad (2.9)$$

The total complexity is summarised in table 2.1 with the complexity for blind operation marked with (*).

Operation	RM	RA
Forming $z[i], \dots, z[i + L]$	$4L, 12L^*$	$3L, 9L^*$
FFT	$2L \log_2(\frac{N_{FFT}}{L}) + 2N_{FFT} \log_2(L)$	$3L \log_2(\frac{N_{FFT}}{L}) + 3N_{FFT} \log_2(L)$
Search	$2N_{FFT}$	$2N_{FFT}$
Interpolation	5	4

Table 2.1: Computational complexity of estimating CFO with the R&B algorithm.

DK The DK estimator has lower complexity, only requiring CORDIC, real additions, and real multiplications. The fixed multiplications by the weights of the Kay estimator can be optimised into a low-energy shift-add circuit, but this is ignored for simplicity. The first-pass CFO correction involves forming the phase corrections of each training symbol as $\phi[i] = \hat{f}_d i$, then applying them with CORDIC.

Operation	RM	RA
Forming $\arg(z[i]), \dots, \arg(z[i + L])$	$2L$	L
Differential Phase	1	$L - 1$
First Pass Correction	$2L$	0
Forming $\arg(z[i]), \dots, \arg(z[i + L])$	$2L$	L
Kay	L	$L - 1$

Table 2.2: Computational complexity of estimating CFO with the DK estimator.

2.2.5 Optimising Bit Width

These frequency recovery algorithms compute an estimate of the normalised CFO f_d/R_s . For n bits of precision, the smallest normalised CFO that can be estimated is 2^{-n} when $0 \leq |f_d/R_s| \leq 1$ is the estimation range. In reality a much smaller range of CFOs needs to be estimated. If the maximum permitted value of $|f_d/R_s| = 2^{-d}$, d bits of precision that would otherwise be redundant can be utilised by scaling each estimate by 2^d during calculation then inserting d bits in the final answer to recover the true normalised CFO. This improves the resolution to $2^{-(n+d)}$ without any increase in energy consumption.

2.3 Phase Recovery

2.3.1 Feed-Forward vs PLL Algorithms

One of the first techniques investigated for phase recovery was a digital phase-locked loop (PLL) to incrementally correct the phase of each symbol with feedback. This dramatically simplifies the phase estimation and correction due to the use of small angle approximations. The CPON spec sets the symbol rate at 30.5 GBd. With an oversampling rate of 8/7 the total samples per second would be 35 GSa/s. This is beyond the achievable clock speed of CMOS, so symbols must be processed in blocks of e.g., 32 symbols with a clock speed of 1.1 GHz. This approach hinders the effectiveness of a PLL, as updates can only be made every 32 symbols. For this reason, feed-forward algorithms that can operate on blocks of data were used instead.

2.3.2 Viterbi & Viterbi

The Viterbi & Viterbi algorithm is well suited to QPSK and removes data phase with the fourth power operation of eq. (2.1). It then applies a maximum-likelihood filter $w_{ML}[i]$ to estimate the phase error due to Wiener phase noise from the combined laser linewidth and any residual CFO as

$$\hat{\theta}_{ML}[i] = \frac{1}{4} \arg \left(\sum_{i=1}^N w_{ML}^*[i] z[i] \right) - \frac{\pi}{4}. \quad (2.10)$$

A phase unwrapper is then applied before the phase error is corrected. This is readily adapted to block-based processing. If blocks of 32 symbols are used again, the ML filter can be applied on all symbols in the block and the estimated phase error applied to the entire block. This is only valid if the phase difference between the start and end of the block is small, meaning the laser linewidth cannot be too large and the preceding frequency recovery algorithm must correct most of the CFO. By increasing the time between phase estimates, block processing also increases the chance of cycle-slips in the phase unwrapper. The CPON spec supplies 1 pilot symbol in every 32-symbol block that can be used to correct this [2]. The correction estimates the number of $\pi/2$ phase jumps n that occur between the unwrapped phase of a data symbol $\hat{\theta}_{PU}$ and its corresponding pilot, then adjusts the unwrapped estimate as $\hat{\theta}'_{PU} = \hat{\theta}_{PU} - n\frac{\pi}{2}$.

Pilots-Only Correction

While Viterbi & Viterbi is commonly used for QPSK, a simpler approach is to compute the phase error from the pilot symbol using eq. (2.2) and apply this to the whole block. This also relies on a small phase change across the block and is more vulnerable to additive noise. However, QPSK allows for $\pm 45^\circ$ error in the phase estimate before an incorrect decision is made, which makes it robust to the drop in accuracy associated with this approach.

2.3.3 Computational Complexity

Viterbi & Viterbi For a block-based approach, the length of the filter will be the block length. Phases are again found with CORDIC. The cost of the phase estimate for each block is shown in table 2.3.

Operation	RM	RA
Forming $z[i], \dots, z[i + N]$	$12N$	$9N$
Computing $\hat{\theta}_{ML}$	$4N + 1$	$3N + 2(N - 1) + 1$

Table 2.3: Computational complexity of block-based Viterbi & Viterbi.

Pilots-Only Only phases are needed, so the entire estimation can be performed with CORDIC. The pilot phase is also known beforehand, so does not contribute to the computational cost. The total cost of the phase estimate is shown in table 2.4.

Operation	RM	RA
Forming $\arg(z[i]) = \hat{\theta}$	1	1

Table 2.4: Computational complexity of Pilots-Only correction.

2.4 Methodology

2.5 Results & Discussion

2.5.1 Maximum CFO

The theoretical limit on the maximum CFO that can be corrected is given by Nyquist as $0.5R_s$. From eq. (2.1) we see that blind operation reduces this limit by a factor of 4. Both the R&B and DK estimators accurately estimate the CFO up to the Nyquist limit as seen in fig. 2.5. Blind operation aliases as expected at $f/R_s = 0.125$.




Figure 2.5: Normalised MSE in CFO estimate (MSE^2/R_s^2) vs Normalised CFO (f/R_s)

2.5.2 Blind vs Data-Aided Operation

The decrease in threshold SNR for longer observation lengths suggests that the performance of the R&B algorithm will improve with blind observation length. This is not the case for the DK estimator, where increasing blind observation length has no effect on threshold SNR. The effect of observation length on FEC SNR is shown for both algorithms in fig. 2.6, where blind R&B shows a clear increase in performance over the blind DK estimator for longer observations and improves on data-aided R&B after ≈ 50 symbols. The DK estimator in blind operation is thus not viable for frequency recovery but blind R&B can provide the best performance at a higher computational cost.




Figure 2.6: FEC SNR vs blind observation length for the R&B and DK estimators. CFO = 2 GHz, laser linewidth = 1 MHz

2.5.3 System Energy

Table 2.5: Minimum combined per-bit energy for the transmitter and carrier recovery for a 10 km passive-split PON (shot-noise limited), in fJ/bit. Values are independent of K .

FR	CR	K	$E_{\text{sys}} \pm \sigma$ (fJ/bit)	FEC SNR (dB)	P_{rx} (dBm)
Diff. Kay	Pilots	1	0.923 ± 0.0050	10.0 ± 0.95	-41.0
		10	0.923 ± 0.0050		-41.0
		100	0.923 ± 0.0050		-41.0
	V&V	1	283 ± 0	15.9 ± 0	-35.1
		10	283 ± 0		-35.1
		100	283 ± 0		-35.1
R&B	Pilots	1	27.2 ± 0.0030	9.39 ± 0.68	-41.7
		10	27.2 ± 0.0030		-41.7
		100	27.2 ± 0.0030		-41.7
	V&V	1	310 ± 0.38	23.4 ± 3.6	-27.7
		10	310 ± 0.38		-27.7
		100	310 ± 0.38		-27.7
R&B blind	Pilots	1	38.7 ± 0.0027	9.59 ± 0.61	-41.5
		10	38.7 ± 0.0027		-41.5
		100	38.7 ± 0.0027		-41.5
	V&V	1	321 ± 0.33	19.2 ± 4.9	-31.9
		10	321 ± 0.33		-31.9
		100	321 ± 0.33		-31.9

Table 2.6: Minimum combined per-bit energy for the transmitter and carrier recovery for a 3×80 km amplified link (ASE + NLI limited), in fJ/bit.

FR	CR	K	$E_{\text{sys}} \pm \sigma$ (fJ/bit)	FEC SNR (dB)	P_{rx} (dBm)
Diff. Kay	Pilots	1	3.34 ± 0.58	10.0 ± 0.95	-18.4
		10	1.15 ± 0.058		-18.4
		100	0.927 ± 0.0058		-18.4
	V&V	1	292 ± 0	15.9 ± 0	-12.5
		10	284 ± 0		-12.5
		100	283 ± 0		-12.5
R&B	Pilots	1	29.3 ± 0.34	9.39 ± 0.68	-19.0
		10	27.4 ± 0.034		-19.0
		100	27.2 ± 0.0034		-19.0
	V&V	1	351 ± 27	21.6 ± 3.0	-6.7
		10	312 ± 1.6		-7.8
		100	310 ± 0.17		-7.7
R&B blind	Pilots	1	40.9 ± 0.32	9.59 ± 0.61	-18.8
		10	38.9 ± 0.032		-18.8
		100	38.7 ± 0.0032		-18.8
	V&V	1	339 ± 13	17.7 ± 3.7	-10.7
		10	323 ± 1.3		-10.7
		100	321 ± 0.14		-10.7

Bibliography

- [1] Cable Television Laboratories, Inc. *Coherent PON Physical Media Dependent Layer 1.0 Specification*. Tech. rep. CPON-SP-PMDv1.0-I01-251216. Issued specification, version I01. CableLabs, Dec. 2025. URL: <https://www.cablelabs.com>.
- [2] Haiquan Cheng et al. “Pilot-symbols-aided cycle slip mitigation for DP-16QAM optical communication systems”. In: *Optics Express* 21.19 (2013), pp. 22166–22172. DOI: 10.1364/OE.21.022166.
- [3] A.N. D’Andrea, Umberto Mengali, and Ruggero Reggiannini. “The Modified Cramer-Rao Bound and Its Applications to Synchronization Problems”. In: *Communications, IEEE Transactions on* 42 (Mar. 1994), pp. 1391–1399. DOI: 10.1109/TCOMM.1994.580247.
- [4] Mark Horowitz. “Computing’s Energy Problem (and what we can do about it)”. In: *2014 IEEE International Solid-State Circuits Conference (ISSCC) Digest of Technical Papers*. 2014, pp. 10–14. DOI: 10.1109/ISSCC.2014.6757323.
- [5] Eric Jacobsen and Peter Kootsookos. “Fast, Accurate Frequency Estimators”. In: *IEEE Signal Processing Magazine* 24.3 (May 2007), pp. 123–125. DOI: 10.1109/MSP.2007.361611.
- [6] Pierluigi Poggiolini et al. “Analytical Modeling of Nonlinear Propagation in Uncompensated Optical Transmission Links”. In: *IEEE Photonics Technology Letters* 23.11 (2011), pp. 742–744. DOI: 10.1109/LPT.2011.2131125.