

Energy-Efficient 200 Gbit/s Transceivers for Optical Access Networks

Joe Leacy

May 18, 2026

Contents

1	Access Networks and Coherent PON	3
1.1	Optical Access Networks	3
1.2	Coherent PON	3
1.3	Modulation and Symbol Rate	4
1.4	Downstream Frame Structure	4
1.5	Key System Tolerances	4
2	Experimental Setup and Methodology	5
2.1	Simulation Framework	5
2.1.1	Network Model	5
2.1.2	Channel Impairments	5
2.1.3	Fixed-Point Arithmetic	6
2.2	Energy Estimation	6
2.3	Receiver Energy	7
2.3.1	Dependence on Word Length	7
2.3.2	Amortisation	7
2.3.3	Total Energy and Power	8
2.4	Transmitter Energy	8
2.4.1	Minimum Required SNR	8
2.4.2	Noise	8
2.5	System Energy	9
3	Equalisation	10
3.1	Introduction	10
3.2	Static Equalisation	10
3.2.1	Equalisation Procedure	10
3.2.2	Filter Length	10
3.2.3	Computational Complexity	11
3.3	Adaptive Equalisation	12
3.3.1	Equalisation Procedure	12
3.3.2	Filter Length	12
3.3.3	Computational Complexity	13
3.3.4	Fixed Point Precision	13
3.4	Methodology	13
4	Carrier Recovery	14
4.1	Introduction	14
4.2	Frequency Recovery	14

4.2.1	Isolating CFO	14
4.2.2	Rife and Boorstyn	15
4.2.3	Differential Phase + Kay Estimator	16
4.2.4	Computational Complexity	18
4.2.5	Optimising Bit Width	19
4.2.6	Maximum CFO	19
4.3	Phase Recovery	20
4.3.1	Feed-Forward vs PLL Algorithms	20
4.3.2	Viterbi & Viterbi	20
4.3.3	Pilots-Only Correction	21
4.3.4	Computational Complexity	21
4.4	Methodology	21
4.4.1	Phase Recovery Comparison	21
4.4.2	System Energy Comparison	22
4.5	Results & Discussion	22
4.5.1	Phase Recovery Performance	22
4.5.2	Blind vs Data-Aided Operation	23
4.5.3	Combined Performance	23
4.6	Conclusions	25

Chapter 1

Access Networks and Coherent PON

1.1 Optical Access Networks

The access network is the final link between an operator's central exchange and the end user, often called the *last mile*. Exponentially growing bandwidth demands from streaming, cloud services, and 5G backhaul are placing severe pressure on this segment.

A passive optical network (PON) uses a point-to-multipoint (P2MP) topology in which a single optical line terminal (OLT) at the exchange serves multiple optical network units (ONUs) at the customer premises through a tree of passive optical splitters, as illustrated in fig. 1.1. No active components are required in the distribution plant, keeping operational costs low.




Figure 1.1: Passive optical network: one OLT serves K ONUs via a passive $1:K$ splitter.

1.2 Coherent PON

Traditional intensity-modulation direct-detection (IM-DD) passive optical networks are approaching their practical limits at data rates approaching 50 Gbps per wavelength. Coherent technology encodes information in the amplitude and phase of the optical signal, allowing modulation formats like QPSK that can carry more bits per symbol. Coherent receivers are also more sensitive, allowing for lower transmit powers. These factors make them a promising alternative to achieve rates of 200 Gbps on an acceptable power budget [4].

The Coherent Passive Optical Network (CPON) specification [1] defines a 100 Gbps single-wavelength P2MP system. This report focuses on the *downstream* direction (OLT to

ONU) where receiver optimisations can be shared by all ONUs; the key physical-layer parameters are described in the following sections.

1.3 Modulation and Symbol Rate

The downstream signal uses non-differential DP-QPSK at a symbol rate of $R_s = 30.5$ GBd. Each polarisation carries 2 bits per symbol by mapping consecutive bit pairs to in-phase and quadrature amplitudes $I, Q \in \{-1, +1\}$, giving QPSK constellation points at $\pm 1 \pm j$. Training and pilot symbols use the same mapping but at amplitude ± 3 , placing them outside the data constellation for unambiguous identification.

1.4 Downstream Frame Structure

The fundamental framing unit is the *DSP subframe* of 3 712 symbols, organised as 116×32 -symbol blocks. Every block begins with a pilot symbol; the first block also carries an 11-symbol training sequence, whose first symbol doubles as pilot index 1.

1.5 Key System Tolerances

The standard supports a 35 dB loss budget from OLT to ONU, which it states will allow a link length of 80 km with a 1:16 splitting ratio or 20 km with a 1:512 splitting ratio. To account for impairments in the core network, an SNR of 35 dB at the OLT must be supported. The OLT transmitter laser must be within ± 1.5 GHz of its DWDM channel centre, and the ONU must acquire signals within ± 1.8 GHz of the grid. In the worst case, the combined transmitter and LO offsets yield a carrier frequency offset (CFO) of up to ≈ 3 GHz that the receiver DSP must correct. Both OLT and ONU lasers must have a linewidth of at most 1 MHz.

The specification requires a pre-FEC bit error ratio of $\leq 2 \times 10^{-2}$, which the soft-decision LDPC FEC reduces to a post-FEC BER of $\leq 10^{-12}$.

Chapter 2

Experimental Setup and Methodology

2.1 Simulation Framework

All DSP algorithms in this report are designed, characterised, and compared through software simulation in MATLAB. The simulated receiver processes DP-QPSK symbols generated at the CPON downstream parameters described in chapter 1.

2.1.1 Network Model

The network under investigation assumes a loss of 35 dB and an SNR of 35 dB at the OLT. The two configurations are 80 km with a 1:16 splitting ratio and 20 km with a 1:512 splitting ratio. The fibre and signal parameters assumed for the network are listed in table 2.1. Nyquist pulse-shaping is assumed.

Table 2.1: System parameters.

Parameter	Value
Central wavelength (λ)	1550 nm
Channel bandwidth (B)	30.5 GHz
Chromatic dispersion coefficient (D)	17 ps/(nm·km)
PMD coefficient (D_{DGD})	0.1 ps/ $\sqrt{\text{km}}$
Wall-plug efficiency (η)	0.1

2.1.2 Channel Impairments

Between the transmitter and receiver, the ideal symbol stream is degraded by the optical channel. Each impairment is applied as a separate stage so that individual DSP blocks can be characterised in isolation or in combination. The relevant models [7] are described below.

Additive White Gaussian Noise AWGN models the combined effect of shot noise and thermal noise. In simulation, the power of the incoming signal is measured and zero-

mean Gaussian noise is added independently to each polarisation at a level set by the target SNR.

Chromatic Dispersion Chromatic dispersion (CD) arises because the group velocity of light in fibre depends on wavelength, so the frequency components of a pulse accumulate a frequency-dependent delay that spreads it over many symbol periods. It is an all-pass effect with a quadratic phase response. It is simulated by transforming the signal to the frequency domain with an FFT, applying the transfer function

$$H(\omega) = \exp\left(j \frac{D\lambda^2}{4\pi c} L \omega^2\right), \quad (2.1)$$

where D is the dispersion coefficient, λ the central wavelength, L the fibre length, and c the speed of light, then transforming back with an IFFT.

Polarisation-Mode Dispersion Polarisation-mode dispersion (PMD) results from random birefringence in the fibre, which causes the two polarisation modes to travel at slightly different speeds and couple into one another along the link. It is simulated as a concatenation of N_{pmd} birefringent sections, each applying a differential group delay of $\pm\tau/2$ to its two principal states in the frequency domain, with a random unitary rotation (the singular vectors of a complex Gaussian matrix) modelling the mode coupling between sections. The mean DGD scales as $\tau \propto D_{\text{PMD}}\sqrt{L}$, with D_{PMD} the fibre PMD coefficient.

Carrier Frequency Offset A carrier frequency offset (CFO) occurs because the transmitter and local-oscillator lasers are free-running and never exactly aligned in centre frequency, as discussed in section 1.5. The resulting beat appears as a linear phase ramp on the received samples. It is simulated by multiplying the signal by $\exp(j 2\pi \Delta f T k)$, where Δf is the frequency offset, T the sample period, and k the sample index.

Phase Noise Phase noise originates from the finite linewidth of the transmitter and LO lasers, whose phases follow a random walk. It is modelled as a Wiener process: independent Gaussian phase increments with variance $\sigma^2 = 2\pi \Delta\nu/R_s$, where $\Delta\nu$ is the combined laser linewidth, are accumulated and applied as a multiplicative phase $\exp(j\phi[k])$.

2.1.3 Fixed-Point Arithmetic

Real hardware DSP implementations operate on fixed-point integers rather than floating-point numbers. Representing a value with n bits limits both the dynamic range and resolution of arithmetic, and the energy cost of addition and multiplication scales with n (see section 2.3). All receiver modules are implemented using MATLAB's `fi` (fixed-point) object, which performs bit-accurate fixed-point arithmetic with configurable overflow and rounding modes.

2.2 Energy Estimation

The energy efficiency of a DSP technique used in a coherent receiver cannot be judged from the receiver alone: a simpler implementation that saves receiver energy typically requires a higher SNR, which the transmitter must compensate for with extra launch

power. A like-for-like comparison between architectures must therefore account for both contributions.

The receiver contribution is modelled from operation counts of real additions and multiplications, with per-operation energies that depend on the fixed-point word length, allowing implementations to be compared without a hardware build. The transmitter contribution is derived from the minimum SNR each implementation requires to meet the pre-FEC BER threshold defined in the CPON standard [1]. The two are combined for a downstream PON to give the per-bit system energy.

2.3 Receiver Energy

The energy consumption of the receiver can be estimated from the number of real multiplications N_M and real additions N_A it performs to process each symbol. The true cost of these operations depends on the hardware implementation. Here the reference values shown in table 2.2 are used to allow for comparison between different DSP algorithms. An accurate estimate of their actual energy consumption would require analysis of their ASIC implementation for a modern process node.

Table 2.2: Energy cost of arithmetic operations for different bit widths in 45 nm CMOS [5].

n (bits)	Add (pJ)	Multiply (pJ)
8	0.03	0.2
32	0.1	3.1

2.3.1 Dependence on Word Length

A straightforward adder implementation performs n bit-wise additions on words of length n , so its energy scales linearly with n . Multiplication accumulates n partial products of length n and so scales as n^2 . These physical scaling laws are imposed *a priori* and the proportionality constants calibrated by least-squares fit through the origin to the two data points of table 2.2, giving

$$E_A(n) = 3.16 n \text{ fJ}, \quad E_M(n) = 3.03 n^2 \text{ fJ}. \quad (2.2)$$

2.3.2 Amortisation

Some DSP algorithms perform computations that are reused over multiple symbols. For example, carrier frequency offset (CFO) varies much slower than the symbol rate so can be estimated less frequently. For a computation that requires N_M multiplications and N_A additions and is used by S symbols, the operations per symbol would be N_M/S and N_A/S .

2.3.3 Total Energy and Power

The energy per symbol for the entire receiver is given by

$$E_{\text{rx},s} = \sum_{k=1}^L 2 [N_{A,k} E_A(n_k) + N_{M,k} E_M(n_k)], \quad (2.3)$$

where the receiver is decomposed into L blocks that use different word lengths n_k according to their required precision, and $N_{A,k}$, $N_{M,k}$ are the addition and multiplication counts per symbol. The factor of 2 is for both polarizations. Blocks operating on the oversampled signal at rate $f > R_s$ have their per-sample operation counts scaled by the oversampling ratio f/R_s . The receiver power is then

$$P_{\text{rx}} = R_s E_{\text{rx},s}, \quad (2.4)$$

and converting to energy per bit allows fair comparison across modulation schemes and symbol rates,

$$E_{\text{rx}} = \frac{E_{\text{rx},s}}{\log_2(M)}, \quad (2.5)$$

where M is the modulation order. Complex multiplications and additions are needed in many DSP algorithms. 4 real multiplications and 2 real additions are assumed for each complex multiplication, and 2 real additions for complex addition.

2.4 Transmitter Energy

The downstream link considered here consists of a single OLT transmitting through a 1: K passive splitter to K ONUs, as shown in fig. 1.1. The OLT power is divided equally between branches, so each ONU receives a $1/K$ share.

2.4.1 Minimum Required SNR

The CPON standard specifies a maximum pre-FEC bit-error rate (BER) of 10^{-2} that must not be exceeded for the forward error-correction (FEC) scheme to function correctly. Each receiver implementation therefore has a minimum required SNR to meet this FEC threshold. This SNR, together with the noise characteristics of the optical network, determines the minimum transmitter power.

2.4.2 Noise

The relationship between SNR and transmitter power P_{tx} depends on the dominant source of noise in the network. Optical access networks typically operate over short lengths at low power, so amplifiers can be omitted and non-linear noise neglected, leaving shot noise as the dominant impairment. The shot noise power at the receiver is $2h\nu B_{\text{eff}}$, where B_{eff} is the bandwidth of the photodetector; this can be matched to the bandwidth of the received signal after the 1: K splitter in ?? with a bandpass filter. The power received by each ONU is $P_{\text{tx}}/K\Gamma$, where Γ is the loss from OLT to ONU. This gives a per-ONU NSR of

$$\text{NSR} = \frac{2K\Gamma h\nu B_{\text{eff}}}{P_{\text{tx}}} + \text{NSR}_0, \quad (2.6)$$

where $\text{NSR} = 1/\text{SNR}$ and NSR_0 is the NSR at the OLT. Rearranging gives the transmitted power as

$$P_{\text{tx}} = \frac{2K\Gamma h\nu B_{\text{eff}}}{\text{NSR} - \text{NSR}_0}. \quad (2.7)$$

The energy per bit transmitted is then

$$E_{\text{tx}} = \frac{P_{\text{tx}}}{KR_s \log_2(M)} = \frac{2\Gamma h\nu}{(\text{NSR} - \text{NSR}_0) \log_2(M)}. \quad (2.8)$$

when $B_{\text{eff}} = R_s$ with Nyquist pulse shaping and bandpass filtering at the receiver.

2.5 System Energy

The total energy consumed per bit for the network is given by

$$E_{\text{sys}} = \frac{1}{\eta} E_{\text{tx}} + E_{\text{rx}}, \quad (2.9)$$

where η is the wall-plug efficiency of the OLT laser.

Chapter 3

Equalisation

3.1 Introduction

3.2 Static Equalisation

3.2.1 Equalisation Procedure

Chromatic dispersion is a static, all-pass impairment and is removed by an equaliser that applies the inverse of the channel response. From the channel model of chapter 2, the compensating transfer function is

$$H_{\text{eq}}(\omega) = \exp\left(-j \frac{D\lambda^2}{4\pi c} L \omega^2\right), \quad (3.1)$$

whose inverse Fourier transform gives the continuous-time impulse response

$$h(t) = \sqrt{\frac{c}{jD\lambda^2 L}} \exp\left(j \frac{\pi c t^2}{D\lambda^2 L}\right). \quad (3.2)$$

In the time domain the equaliser is a single FIR filter per polarisation whose taps are eq. (3.2) sampled at the receiver period T and truncated to N_{CD} coefficients. Alternatively, eq. (3.1) can be applied directly in the frequency domain. As the data stream is continuous, the filter is applied block-by-block by overlap-save: each block of N samples is transformed with an FFT, multiplied by the precomputed samples of H_{eq} , and inverse transformed, discarding the first $N_{\text{CD}} - 1$ samples of every block to remove the circular-convolution wrap-around. This leaves $N - N_{\text{CD}} + 1$ valid output samples per block.

3.2.2 Filter Length

The equaliser must span the channel memory introduced by dispersion. The pulse spread across the signal bandwidth is $\Delta\tau = D \Delta\lambda L$, so the number of taps required is

$$N_{\text{CD}} = \frac{\Delta\tau}{T} + 1, \quad (3.3)$$

where T is the receiver sampling period. With an oversampling rate of 2, table 3.1 lists the taps needed for the two networks of chapter 2.

Table 3.1: Taps to compensate chromatic dispersion for both network models.

L (km)	$\Delta\tau$ (ps)	N_{CD}
20	83	6
80	332	21

3.2.3 Computational Complexity

A radix-2 FFT of size N decomposes into $\log_2 N$ stages, each containing $N/2$ butterflies. Each butterfly applies one twiddle-factor complex multiplication followed by a complex addition and a complex subtraction. With 4 real multiplications and 2 real additions per complex multiplication, and 2 real additions per complex addition, a butterfly costs 4 real multiplications and $2 + 2 \times 2 = 6$ real additions. The full transform therefore requires

$$N_M^{\text{FFT}} = 2N \log_2 N, \quad N_A^{\text{FFT}} = 3N \log_2 N. \quad (3.4)$$

Each overlap-save block needs a forward FFT, an inverse FFT of equal cost, and N complex multiplications by the precomputed filter spectrum (the filter FFT is computed once offline as the channel is static). This is $2N_M^{\text{FFT}} + 4N$ real multiplications and $2N_A^{\text{FFT}} + 2N$ real additions per block, shared over the $N - N_{\text{CD}} + 1$ valid outputs. Scaling by the oversampling ratio $S = f/R_s$ to express the cost per symbol gives

$$\begin{aligned} N_{M,\text{OS}} &= S \frac{4N \log_2 N + 4N}{N - N_{\text{CD}} + 1}, \\ N_{A,\text{OS}} &= S \frac{6N \log_2 N + 2N}{N - N_{\text{CD}} + 1}. \end{aligned} \quad (3.5)$$

In the time domain, direct convolution with the N_{CD} -tap complex filter costs N_{CD} complex multiplications and $N_{\text{CD}} - 1$ complex additions per output sample. Converting to real operations and scaling by S , the per-symbol cost is

$$N_{M,\text{TD}} = 4SN_{\text{CD}}, \quad N_{A,\text{TD}} = S(4N_{\text{CD}} - 2). \quad (3.6)$$

The FFT can be simplified by rounding all twiddle-factors to a power of 2. This eliminates all real multiplications in performing the FFT/IFFT, replacing them with bit-shifts of negligible cost.

Operation	Time-domain	Overlap-save	Overlap-save (2^k twiddle)
Real multiplications	$4SN_{\text{CD}}$	$S \frac{4N \log_2 N + 4N}{N - N_{\text{CD}} + 1}$	$S \frac{4N}{N - N_{\text{CD}} + 1}$
Real additions	$S(4N_{\text{CD}} - 2)$	$S \frac{6N \log_2 N + 2N}{N - N_{\text{CD}} + 1}$	$S \frac{6N \log_2 N + 2N}{N - N_{\text{CD}} + 1}$

Table 3.2: Per-symbol computational cost of the time-domain and overlap-save chromatic dispersion equalisers, including overlap-save with twiddle factors rounded to a power of 2.

Evaluating these expressions with the N_{CD} of table 3.1 gives the per-symbol operation counts for each network in table 3.3.

	Time-domain		Overlap-save		Overlap-save (2^k twiddle)	
L (km)	RM	RA	RM	RA	RM	RA
20						
80						

Table 3.3: Real multiplications (RM) and real additions (RA) per symbol for each chromatic dispersion equaliser implementation, evaluated with the N_{CD} and N of table 3.1.

3.3 Adaptive Equalisation

3.3.1 Equalisation Procedure

Tracking PMD is done with a 2x2 MIMO equaliser that employs 4 finite impulse response (FIR) filters to give outputs

$$x_{\text{out}} = \mathbf{h}_{xx}^H \mathbf{x}_{\text{in}} + \mathbf{h}_{xy}^H \mathbf{y}_{\text{in}}, \quad y_{\text{out}} = \mathbf{h}_{yx}^H \mathbf{x}_{\text{in}} + \mathbf{h}_{yy}^H \mathbf{y}_{\text{in}}. \quad (3.7)$$

The filter weights are updated using an error estimate as

$$\begin{aligned} \mathbf{h}_{xx} &\leftarrow \mathbf{h}_{xx} + \mu e_x \mathbf{x}_{\text{in}} x_{\text{out}}^*, & \mathbf{h}_{xy} &\leftarrow \mathbf{h}_{xy} + \mu e_x \mathbf{y}_{\text{in}} x_{\text{out}}^*, \\ \mathbf{h}_{yx} &\leftarrow \mathbf{h}_{yx} + \mu e_y \mathbf{x}_{\text{in}} y_{\text{out}}^*, & \mathbf{h}_{yy} &\leftarrow \mathbf{h}_{yy} + \mu e_y \mathbf{y}_{\text{in}} y_{\text{out}}^*. \end{aligned} \quad (3.8)$$

where the convergence parameter μ is chosen to be a power of 2 for hardware simplicity. The error terms are produced blindly using the constant modulus algorithm (CMA) [coherent receivers], which exploits the fact that an ideal QPSK constellation lies on a circle of fixed radius. CMA minimises the cost function $J = \mathbb{E}[(|x_{\text{out}}|^2 - R)^2]$, where the dispersion constant $R = \mathbb{E}[|s|^4]/\mathbb{E}[|s|^2]$ is set by the transmitted symbols s and reduces to a constant for QPSK. Taking the stochastic gradient of J yields the error estimates

$$e_x = (R - |x_{\text{out}}|^2), \quad e_y = (R - |y_{\text{out}}|^2). \quad (3.9)$$

The update step from eq. (3.8) can be simplified by only considering the signs of the error and output terms with no performance penalty [Cardenas:14]. The output is a complex number so a complex signum function is defined as

$$csgn(z) = sgn(\text{Re}(z)) + jsngn(\text{Im}(z)). \quad (3.10)$$

3.3.2 Filter Length

There must be sufficient taps in the filter to cover the channel memory introduced by PMD. The mean differential group delay (DGD) $\langle \Delta\tau \rangle = D\sqrt{L}$ where D is the DGD parameter and L is the link length. The instantaneous DGD is modelled by a Maxwellian distribution [Mantzoukis:11]

$$\mathbb{P}[\Delta\tau > x] \approx \frac{4}{\pi} \frac{x}{\langle \Delta\tau \rangle} \exp\left(-\frac{4x^2}{\pi \langle \Delta\tau \rangle^2}\right). \quad (3.11)$$

The number of taps can be found for a given outage probability from eq. (3.11) as

$$N = \frac{x}{T} + 1 \quad (3.12)$$

where T is the sampling period of the receiver. For an access network where $L \leq 80$ km, and a typical DGD parameter $D = 0.1$ ps/ $\sqrt{\text{km}}$ then $\langle \Delta\tau \rangle \leq 0.9$ ps. With an oversampling rate of 2, the sampling period $T = 16.4$ ps. Given that an outage probability of 10^{-12} only requires $x = 4.81\langle \Delta\tau \rangle = 4.3$ ps, few taps will be needed for reliable performance.

3.3.3 Computational Complexity

The per-sample cost of the weight updates is compared below for the direct and sign-sign implementations with N taps.

Operation	Direct	Sign-Sign
Real multiplications	$16N + 4$	0
Real additions	$16N$	$16N$

Table 3.4: Per-symbol computational cost of the direct and sign-sign weight updates.

The adaptive equaliser must be parallelised to support a 30.5 GBd symbol rate. The input samples are split into P lanes, with the input vector of lane p

$$\mathbf{x}_p = [\mathbf{x}[mP + p], \mathbf{x}[mP + p - 1], \dots, \mathbf{x}[mP + p - (N - 1)]] . \quad (3.13)$$

The update step for the shared weights is given by

$$\Delta \mathbf{h} = \mu \sum_{p=0}^{P-1} e[p] \mathbf{x}_p x_{\text{out}}[p]. \quad (3.14)$$

A natural choice for blocks of 32 symbols is $P = 32$, allowing the pilot symbols with a larger modulus to be processed separately. With an oversampling rate of 2, each lane would need to operate with a sampling frequency of $2 \cdot 30.5/32 = 1.9$ GHz which is reasonable for modern CMOS. This scales the number of operations in table 3.4 by P .

3.3.4 Fixed Point Precision

The convergence parameter μ in eq. (3.8) is typically $\approx 10^{-3}$ which would require at least 10 bits of precision to represent. Computing the update term to this precision adds unnecessary computational cost, so it should instead be computed at lower precision and bit-shifted (equivalent to setting μ to a power of 2) to the least significant bits of the weights before updating. This can lead to the most significant bits being redundant. These bits waste energy when the weights change sign and when forming the outputs in eq. (3.7) if the weights are negative. Jiang et al. [Jiang;26] propose removing these redundant bits with an optimisation over the precision of each weight and demonstrate a decrease in energy consumption with a slight performance loss. The decrease in energy consumption can be estimated directly from the frequency of sign changes in the weights and the proportion of multiplications in eq. (3.7) that use negative weights. For d redundant bits, a sign change is equivalent to the cost of a d -bit addition and the additional cost in forming the outputs will be a d -bit multiplication.

3.4 Methodology

Chapter 4

Carrier Recovery

4.1 Introduction

A coherent receiver recovers the transmitted symbols by mixing the incoming optical signal with a local oscillator (LO) laser, allowing the amplitude and phase of each symbol to be extracted directly. However, the transmitter and LO lasers will in general have different centre frequencies and independent phase trajectories, producing a carrier frequency offset (CFO) and accumulated phase noise on the received symbols. Carrier recovery is the DSP stage responsible for estimating and removing these impairments before symbol decisions are made.

Carrier recovery is divided into two sequential stages. Frequency recovery performs a coarse estimate of the CFO and corrects it so that only slow residual phase variations remain. Phase recovery then tracks and removes these residual phase errors on a block-by-block basis. The algorithm chosen for each stage determines both the minimum achievable SNR and the receiver DSP energy, making algorithm selection central to the system-level energy comparison of chapter 2.

This chapter describes the Rife and Boorstyn (R&B) and Differential Kay (DK) frequency recovery algorithms and the Viterbi & Viterbi (V&V) and pilots-only phase recovery algorithms, analysing their estimation accuracy and computational complexity. The optimal fixed-point word length for each combination is then found, and the resulting system energies are compared across shot-noise and amplified link scenarios.

4.2 Frequency Recovery

4.2.1 Isolating CFO

All frequency recovery algorithms first require that the phase of the data be removed before estimation. For QPSK symbols $c[i] = e^{j\frac{\pi}{4} + m[i]\frac{\pi}{2}}$, this can be done by raising each received symbol $x[i] = c[i]e^{j(2\pi f_d i + \theta[i])} + n[i]$ to the fourth power to give

$$z[i] = x^4[i] = e^{j(8\pi f_d i + 4\theta[i])} + n'[i] \quad (4.1)$$

where θ is phase noise, f_d is the CFO, and n is additive noise. This allows all data symbols to be used in CFO estimation, but amplifies phase and additive noise, which

degrades the estimation quality. The fourth power operation also requires three complex multiplications. An alternative is to use pilot symbols where the data phase is known beforehand, giving

$$z[i] = x[i]c^*[i] = e^{j(2\pi f_d i + \theta[i])} + n[i], \quad (4.2)$$

which is more robust to noise and less computationally expensive. Some algorithms only require $\arg(z[i])$, which is efficiently found using CORDIC as

$$\arg(z[i]) = \{4 \arg(x[i])\}_{-\pi}^{\pi} \quad (4.3)$$

for blind operation or

$$\arg(z[i]) = \{\arg(x[i]) - \arg(c[i])\}_{-\pi}^{\pi} \quad (4.4)$$

for data-aided operation, where the wrap operation can be done with no additional cost by aligning the wrap range to the overflow range of the fixed-point value used to store the phase.

4.2.2 Rife and Boorstyn

The value of the CFO will appear as the strongest peak in the frequency spectrum of $z[i], z[i+1], \dots, z[i+L]$ where L is the observation length. The Rife and Boorstyn (R&B) algorithm first performs an FFT on the observed sequence and finds the index k corresponding to the strongest peak. This coarse estimation is refined with a parabolic interpolation on the complex FFT values X_{k-1}, X_k, X_{k+1} [6] to give

$$\delta = -\Re \left[\frac{X_{k+1} - X_{k-1}}{2X_k - X_{k-1} - X_{k+1}} \right]. \quad (4.5)$$

The updated index $k' = k + \delta$ is then scaled by the sampling rate to find the true CFO. The resolution in k is $\epsilon(k) = 1/2N_{FFT}$. The CPON specification provides only $L = 11$ training symbols, so data-aided operation without zero-padding gives $\epsilon(\hat{f}_d) \approx 0.05 R_s$ — too coarse for the downstream phase recovery to correct. Zero-padding the training sequence to a larger N_{FFT} improves the resolution without requiring additional training data. The normalised mean squared-error (MSE) in the estimated CFO is shown in fig. 4.1 for data-aided operation with 11 training symbols. A CFO estimate is computed for both polarizations and averaged to reduce the MSE. The modified Cramer-Rao bound (MCRB) [3] provides a lower bound on the achievable MSE and is given by

$$\text{MCRB} = \frac{3}{2\pi^2 L^3 \times \text{SNR}}. \quad (4.6)$$

R&B achieves this bound above a threshold SNR of ≈ 4 dB, where the threshold SNR is defined as the SNR below which the estimator begins to deviate from the MCRB. Below this, the noise has a significant contribution to the FFT values and can lead to an incorrect index being chosen.




Figure 4.1: Normalised MSE of the R&B algorithm with FFT size $N_{FFT} = 512$ compared to the MCRB for $L = 11$ training symbols.

Data-aided operation is constrained by the length of the training sequence, but the accuracy of the R&B algorithm improves with longer observation lengths, as shown in fig. 4.2 for non-data-aided operation. The SNR threshold at which this algorithm begins to deviate from the MCRB is initially higher with data-aided operation due to the amplification of noise, but decreases for longer observation lengths.




Figure 4.2: Normalised MSE of the R&B algorithm for non-data-aided operation and different observation lengths.

4.2.3 Differential Phase + Kay Estimator

The R&B algorithm exhibits good performance even at low SNRs, but is computationally expensive. The FFT requires $\frac{N_{FFT}}{2} \log_2(N_{FFT})$ complex multiplications and

$N_{FFT} \log_2(N_{FFT})$ complex additions. The fine search from eq. (4.5) also requires a complex division. A cheaper approach would involve computing the phase difference

$$\arg(z[i]) - \arg(z[i-1]) = 2\pi f_d T_s + \theta[i] - \theta[i-1] + \eta[i] \quad (4.7)$$

over several symbols and averaging to remove the effect of θ . Estimating f_d in this manner is vulnerable to the effect of additive noise on the phase $\eta[i]$, so exhibits poor accuracy. This can be improved with a second pass using the Kay estimator. The Kay estimator uses a least-squares approach, computing the estimated CFO as

$$\hat{f}_d = \frac{R_s}{2\pi} \sum_{k=1}^L \frac{6k(L-k)}{L(L^2-1)} \arg(z[k]z^*[k-1]). \quad (4.8)$$

where $\arg(z[k]z^*[k-1])$ computes the phase difference from eq. (4.7) and $\frac{6k(L-k)}{L(L^2-1)}$ is a weighting term that prioritises the contribution of phase differences in the middle of the observed sequence. This estimator is able to achieve the MCRB, but the least-squares approach considers phase differences of non-adjacent symbols, making it vulnerable to phase wrapping that distorts the estimate when the CFO is high enough. When the total phase drift $\Delta\phi = 2\pi f_d L T_s$ across the observation window exceeds 2π , a wrap shifts the estimated CFO by $\pm 1/(L T_s)$. A first-pass coarse CFO correction must therefore be applied before the Kay estimator to keep the residual CFO within the unambiguous range. The performance of the Differential Kay (DK) estimator for data-aided operation is shown in fig. 4.3. The MCRB is achieved above a threshold SNR of ≈ 7 dB. This is higher than for R&B, but at a reduced computational cost. Unlike the R&B algorithm, the threshold SNR of the DK estimator does not depend on observation length, as shown in fig. 4.4.




Figure 4.3: Normalised MSE of frequency estimation with the DK estimator for $L = 11$ training symbols.




Figure 4.4: Normalised MSE of the DK estimator with blind operation for different observation lengths.

4.2.4 Computational Complexity

Here, the computational complexity of both frequency recovery algorithms is compared in terms of real multiplications and real additions. CORDIC is used to find phases and is assumed to cost the same as a real multiplication. Real division is also assumed to cost the same as multiplication since it is only performed once (this would be a suitable approximation if look-up tables are used to find reciprocals at some cost to accuracy). All analysis is for data-aided operation, avoiding the fourth power operation in eq. (4.1), and applies to a single polarization. If averaging is done over both polarizations, then all operations are doubled.

R&B The R&B algorithm requires an FFT followed by N_{FFT} comparisons of the magnitudes in the coarse search. Many of the FFT entries are 0 due to padding, which reduces the total number of operations needed. The interpolation requires the real part of a complex division

$$\Re \left[\frac{a + bj}{c + dj} \right] = \frac{ac + bd}{c^2 + d^2}. \quad (4.9)$$

The total complexity is summarised in table 4.1, with the complexity for blind operation marked with (*).

Operation	RM	RA
Forming $z[i], \dots, z[i + L]$	$4L, 12L^*$	$3L, 9L^*$
FFT	$2L \log_2(\frac{N_{FFT}}{L}) + 2N_{FFT} \log_2(L)$	$3L \log_2(\frac{N_{FFT}}{L}) + 3N_{FFT} \log_2(L)$
Search	$2N_{FFT}$	$2N_{FFT}$
Interpolation	5	4

Table 4.1: Computational complexity of estimating CFO with the R&B algorithm.

DK The DK estimator has lower complexity, only requiring CORDIC, real additions, and real multiplications. The fixed multiplications by the weights of the Kay estimator can be optimised into a low-energy shift-add circuit, but this is ignored for simplicity. The first-pass CFO correction involves forming the phase corrections of each training symbol as $\phi[i] = \hat{f}_d i$, then applying them with CORDIC.

Operation	RM	RA
Forming $\arg(z[i]), \dots, \arg(z[i + L])$	$2L$	L
Differential Phase	1	$L - 1$
First Pass Correction	$2L$	0
Forming $\arg(z[i]), \dots, \arg(z[i + L])$	$2L$	L
Kay	L	$L - 1$

Table 4.2: Computational complexity of estimating CFO with the DK estimator.

4.2.5 Optimising Bit Width

These frequency recovery algorithms compute an estimate of the normalised CFO f_d/R_s . For n bits of precision, the smallest normalised CFO that can be estimated is 2^{-n} for an estimation range $0 \leq |f_d/R_s| \leq 1$. In reality, a much smaller range of CFOs needs to be estimated. If the maximum permitted value of $|f_d/R_s| = 2^{-d}$, d bits of precision that would otherwise be redundant can be utilised by scaling each estimate by 2^d during calculation, then inserting d bits in the final answer to recover the true normalised CFO. This improves the resolution to $2^{-(n+d)}$ without any increase in energy consumption.

4.2.6 Maximum CFO

The theoretical limit on the maximum CFO that can be corrected is given by Nyquist as $0.5R_s$. From eq. (4.1), we see that blind operation reduces this limit by a factor of 4. Both the R&B and DK estimators accurately estimate the CFO up to the Nyquist limit, as seen in fig. 4.5. Blind operation aliases as expected at $f/R_s = 0.125$, making it viable for CFOs up to 3.8 GHz with a symbol rate of 30.5 GHz.




Figure 4.5: Normalised MSE in CFO estimate (MSE^2/R_s^2) vs Normalised CFO (f/R_s)

4.3 Phase Recovery

4.3.1 Feed-Forward vs PLL Algorithms

One of the first techniques investigated for phase recovery was a digital phase-locked loop (PLL) to incrementally correct the phase of each symbol with feedback. This dramatically simplifies the phase estimation and correction due to the use of small-angle approximations. The CPON specification sets the symbol rate at 30.5 GBd. With an oversampling rate of 8/7, the total samples per second would be 35 GSa/s. This is beyond the achievable clock speed of CMOS, so symbols must be processed in blocks of e.g., 32 symbols with a clock speed of 1.1 GHz. This approach hinders the effectiveness of a PLL, as updates can only be made every 32 symbols. For this reason, feed-forward algorithms that can operate on blocks of data were used instead.

4.3.2 Viterbi & Viterbi

The Viterbi & Viterbi algorithm is well suited to QPSK and removes data phase with the fourth power operation of eq. (4.1). It then applies a maximum-likelihood filter $w_{ML}[i]$ to estimate the phase error due to Wiener phase noise from the combined laser linewidth and any residual CFO as

$$\hat{\theta}_{ML}[i] = \frac{1}{4} \arg \left(\sum_{k=1}^N w_{ML}^*[k] z[k] \right) - \frac{\pi}{4}. \quad (4.10)$$

A phase unwrapper is then applied before the phase error is corrected. This is readily adapted to block-based processing. If blocks of 32 symbols are used again, the ML filter can be applied to all symbols in the block and the estimated phase error applied to the entire block. This is only valid if the phase difference between the start and end of the block is small, meaning the laser linewidth cannot be too large and the preceding frequency recovery algorithm must correct most of the CFO. By increasing the time between phase

estimates, block processing also increases the chance of cycle-slips in the phase unwrapper. The CPON specification provides 1 pilot symbol in every 32-symbol block that can be used to correct this [2]. The correction estimates the number of $\pi/2$ phase jumps n that occur between the unwrapped phase of a data symbol $\hat{\theta}_{PU}$ and its corresponding pilot, then adjusts the unwrapped estimate as $\hat{\theta}'_{PU} = \hat{\theta}_{PU} - n\frac{\pi}{2}$.

4.3.3 Pilots-Only Correction

While Viterbi & Viterbi is commonly used for QPSK, a simpler approach is to compute the phase error from the pilot symbol using eq. (4.2) and apply this to the whole block. This also relies on a small phase change across the block and is more vulnerable to additive noise. However, QPSK allows for $\pm 45^\circ$ error in the phase estimate before an incorrect decision is made, which makes it robust to the drop in accuracy associated with this approach.

4.3.4 Computational Complexity

Viterbi & Viterbi For a block-based approach, the length of the filter will be the block length. Phases are again found with CORDIC. The cost of the phase estimate for each block is shown in table 4.3.

Operation	RM	RA
Forming $z[i], \dots, z[i + N]$	$12N$	$9N$
Computing $\hat{\theta}_{ML}$	$4N + 1$	$3N + 2(N - 1) + 1$

Table 4.3: Computational complexity of block-based Viterbi & Viterbi.

Pilots-Only Only phases are needed, so the entire estimation can be performed with CORDIC. The pilot phase is also known beforehand, so does not contribute to the computational cost. The total cost of the phase estimate is shown in table 4.4.

Operation	RM	RA
Forming $\arg(z[i]) = \hat{\theta}$	1	1

Table 4.4: Computational complexity of Pilots-Only correction.

4.4 Methodology

4.4.1 Phase Recovery Comparison

The two phase recovery algorithms are compared in isolation by simulating an AWGN channel with additive Wiener phase noise, without chromatic dispersion, PMD, or fibre nonlinearity, so that any performance difference is attributable solely to the phase estimation stage. BER is estimated by Monte Carlo simulation over independent realisations for each combination of laser linewidth and SNR, with the phase ambiguity inherent to QPSK resolved by exhaustive search over the four $\pi/2$ rotations before symbol decisions are made.

4.4.2 System Energy Comparison

The energy framework of chapter 2 is applied to compare all combinations of frequency and phase recovery algorithm. All simulations target the CPON specification.

For each algorithm combination, the fixed-point word lengths of the frequency recovery (n_{FR}) and phase recovery (n_{PR}) stages are varied independently over a grid. At each grid point, Monte Carlo simulation over multiple independent realisations estimates the FEC SNR which determines the required transmitter power via eq. (2.6). The receiver DSP energy at that precision is computed from the operation counts of the preceding sections via eq. (2.5), and the transmitter energy follows from eq. (2.8). The precision combination that minimises the mean system energy of eq. (2.9) is then selected for each algorithm pair and splitting ratio.

The reported energies should be interpreted as a *comparative* metric. The 45 nm CMOS reference values of table 2.2 provide a consistent basis for ranking implementations, but do not constitute absolute power predictions for a real device. Additionally, the energy cost of applying the estimated corrections to each received symbol is excluded, as this contribution is implementation-independent and does not affect relative comparisons.

4.5 Results & Discussion

4.5.1 Phase Recovery Performance

To replace Viterbi & Viterbi with the simplified pilots-only phase correction, it must exhibit comparable performance over the range of linewidths that could be expected during operation of the network. The BER for each phase recovery algorithm as SNR varies is shown in fig. 4.6. Both algorithms have almost identical performance and are reliable up to 10 MHz, well beyond the maximum specified in the CPON spec.




Figure 4.6: Performance comparison of V&V and Pilots-Only phase correction

4.5.2 Blind vs Data-Aided Operation

The decrease in threshold SNR for longer observation lengths suggests that the performance of the R&B algorithm will improve with blind observation length D . This is not the case for the DK estimator, where increasing blind observation length has no effect on threshold SNR. The effect of observation length on FEC SNR is shown for both algorithms in fig. 4.7, where blind R&B shows a clear increase in performance over the blind DK estimator for longer observations and improves on data-aided R&B after ≈ 50 symbols. The DK estimator in blind operation is thus not viable for frequency recovery, but blind R&B can provide the best performance at a higher computational cost.




Figure 4.7: FEC SNR vs blind observation length for the R&B and DK estimators. CFO = 2 GHz, laser linewidth = 1 MHz.

4.5.3 Combined Performance

The best performance achieved by each carrier recovery combination is shown in table 4.5. The most computationally expensive combination of blind R&B and Viterbi & Viterbi achieves the lowest FEC SNR. DK and data-aided R&B exhibit almost identical best-case performance.

Table 4.5: Minimum mean FEC SNR achievable by each frequency/phase recovery pair over the fixed-point precision grid (energy cost disregarded), with the precision combo at which the minimum occurs.

FR	PR	Min FEC SNR $\pm \sigma$ (dB)	(n_{FR}, n_{PR})
DK	Pilots	7.94 ± 0.40	(14, 8)
DK	V&V	7.64 ± 0.25	(14, 8)
R&B	Pilots	7.93 ± 0.40	(12, 14)
R&B	V&V	7.64 ± 0.27	(10, 8)
R&B blind (D=512)	Pilots	7.50 ± 0.08	(12, 10)
R&B blind (D=256)	V&V	7.39 ± 0.08	(14, 10)

The best-case performance is not of interest in an energy-efficient design. Instead we wish to find the implementation that uses the least system energy to achieve acceptable performance. Figure 4.8 shows the minimum combined energy required to reach the pre-FEC threshold for each frequency/phase recovery pair, with the energy-optimal fixed-point precision and the corresponding energy breakdown listed in table 4.6. Amortisation means phase recovery dominates the receiver energy. The high complexity of Viterbi & Viterbi means that the optimal implementation uses few bits of precision to reduce the receiver energy at the cost of a high FEC SNR. Viterbi & Viterbi evidently suffers more from this reduced precision than pilots-only as quantisation error accumulates over multiple operations, resulting in greater energy consumption by the transmitter and receiver. The choice of frequency recovery algorithm is less significant, with DK saving ≈ 30 fJ/bit over R&B when a CFO estimate is applied to each subframe (3712 symbols).




Figure 4.8: Energy breakdown for carrier recovery. Left: (20 km, 1:512). Right: (80 km, 1:16).

Table 4.6: Energy-optimal fixed-point configuration for each frequency/phase recovery pair under the 35 dB flat-loss shot-noise model. Energies in fJ/bit; results are independent of the splitting ratio K .

FR	PR	$(n_{\text{FR}}, n_{\text{PR}})$	$E_{\text{tx}} / E_{\text{rx}}$	$E_{\text{sys}} \pm \sigma$	FEC SNR $\pm \sigma$ (dB)
DK	Pilots	(4, 4)	26.5 / 3.07	29.5 ± 3.4	8.11 ± 0.50
DK	V&V	(6, 4)	225 / 956	1181 ± 40	17.2 ± 1.3
R&B	Pilots	(2, 4)	32.6 / 28.5	61.1 ± 4.3	9.01 ± 0.54
R&B	V&V	(2, 4)	238 / 980	1218 ± 25	17.6 ± 0.71
R&B blind ($D=32$)	Pilots	(2, 4)	33.0 / 40.0	73.0 ± 4.0	9.06 ± 0.51
R&B blind ($D=64$)	V&V	(2, 2)	748 / 329	1077 ± 987	19.7 ± 4.6

If the CFO remains constant over a longer period, the energy contribution from CFO estimation becomes negligible. With the frequency recovery energy amortised away, the energy-optimal pilots-only configurations are listed in table 4.7. Blind R&B now becomes the most efficient frequency recovery algorithm, since its high observation length lowers the FEC SNR at no per-symbol cost, narrowly beating DK.

Table 4.7: Energy-optimal configuration when the frequency recovery energy is amortised away (CFO constant over many symbols). Energies in fJ/bit; results are independent of the splitting ratio K .

FR	PR	$(n_{\text{FR}}, n_{\text{PR}})$	$E_{\text{tx}} / E_{\text{rx}}$	$E_{\text{sys}} \pm \sigma$	FEC SNR $\pm \sigma$ (dB)
DK	Pilots	(8, 4)	25.8 / 1.91	27.7 ± 2.8	8.01 ± 0.44
R&B blind ($D=256$)	Pilots	(12, 4)	23.1 / 1.91	25.0 ± 0.39	7.55 ± 0.07

4.6 Conclusions

The robustness to both additive and phase noise of QPSK allows the phase recovery block of a coherent receiver to be dramatically simplified by just using pilot symbols with no noticeable drop in performance, resulting in considerable energy savings. Frequency recovery can also be implemented with simple algorithms to save energy; however, the efficacy of doing so depends on how often the CFO must be estimated. For access networks dominated by shot noise where dispersive effects have been compensated, carrier recovery performance is independent of the length and splitting ratio when constrained to a given loss from OLT to ONU.

Bibliography

- [1] Cable Television Laboratories, Inc. *Coherent PON Physical Media Dependent Layer 1.0 Specification*. Tech. rep. CPON-SP-PMDv1.0-I01-251216. Issued specification, version I01. CableLabs, Dec. 2025. URL: <https://www.cablelabs.com>.
- [2] Haiquan Cheng et al. “Pilot-symbols-aided cycle slip mitigation for DP-16QAM optical communication systems”. In: *Optics Express* 21.19 (2013), pp. 22166–22172. DOI: 10.1364/OE.21.022166.
- [3] A.N. D’Andrea, Umberto Mengali, and Ruggero Reggiannini. “The Modified Cramer-Rao Bound and Its Applications to Synchronization Problems”. In: *Communications, IEEE Transactions on* 42 (Mar. 1994), pp. 1391–1399. DOI: 10.1109/TCOMM.1994.580247.
- [4] Md Saifuddin Faruk et al. “Coherent Passive Optical Networks: Why, When, and How”. In: *IEEE Communications Magazine* 59.12 (2021), pp. 112–117. DOI: 10.1109/MCOM.010.2100503.
- [5] Mark Horowitz. “Computing’s Energy Problem (and what we can do about it)”. In: *2014 IEEE International Solid-State Circuits Conference (ISSCC) Digest of Technical Papers*. 2014, pp. 10–14. DOI: 10.1109/ISSCC.2014.6757323.
- [6] Eric Jacobsen and Peter Kootsookos. “Fast, Accurate Frequency Estimators”. In: *IEEE Signal Processing Magazine* 24.3 (May 2007), pp. 123–125. DOI: 10.1109/MSP.2007.361611.
- [7] Darli Mello and Fabio Barbosa. *Digital Coherent Optical Systems: Architecture and Algorithms*. Jan. 2021. ISBN: 978-3-030-66540-1. DOI: 10.1007/978-3-030-66541-8.